

PATHI MANIKANTA

AI Evaluation Specialist | ML Engineer | RLHF/RLAIF

pmanikanta818@gmail.com | Kadapa, India | [LinkedIn](#) | github.com/PSPMANI | pspmeni.github.io

PROFESSIONAL SUMMARY

AI Evaluation Specialist and Machine Learning Engineer with hands-on experience building and deploying end-to-end ML pipelines, evaluating frontier AI systems, and contributing to RLHF/RLAIF workflows. Skilled in Python, scikit-learn, TensorFlow/Keras, NLP, computer vision, model deployment, and agentic tool-use evaluation. Earned promotion from attempter to reviewer at Scale AI (Outlier) for maintaining 95%+ grading consistency across 100+ reviews. Built and deployed evaluation and ML systems, including a live agentic-trajectory auditor and a classifier reaching 92.5% accuracy with proper data leakage detection and GroupKFold cross-validation.

TECHNICAL SKILLS

Programming Languages: Python, Java, C/C++, JavaScript, SQL, HTML, CSS, Embedded C

ML & AI Frameworks: scikit-learn, TensorFlow, Keras, pandas, NumPy, NLTK, OpenCV, SHAP, Streamlit, matplotlib, seaborn

AI Evaluation: LLM evaluation, RLHF, RLAIF, rubric design, agentic trajectory review, tool-use evaluation, API evaluation, red-teaming, prompt engineering, pytest verifiers, fact-checking, data annotation

ML Techniques: Classification, regression, model training & evaluation, feature engineering, data preprocessing, signal processing, cross-validation, GroupKFold, hyperparameter tuning, data leakage detection, model deployment, model explainability (SHAP), NLP, TF-IDF, text classification, computer vision, face detection & recognition

Tools & Platforms: Git, GitHub, MATLAB, Excel, Power BI, Arduino, AWS (EC2, S3, IAM, VPC), Google Workspace, VS Code, Jupyter, pytest

Web & Frameworks: React, Node.js, REST APIs, JSON, HTML/CSS

Spoken Languages: English, Telugu, Hindi, Tamil

PROFESSIONAL EXPERIENCE

AI Evaluation Specialist, Freelance | Scale AI (Outlier) | 2026 - Present

- Promoted from attempter to reviewer within 3 months after designing single-turn agentic evaluation tasks with consistent grading accuracy and full guideline compliance.
- Authored 50+ binary pass/fail rubrics (15-25 criteria each) and developed pytest verifiers that programmatically validate task outputs using systematic test design methodology.
- Audited 100+ agentic tool-use trajectories, identifying failure patterns in call sequencing, argument accuracy, and instruction-following to improve frontier model performance.
- Architected multimodal evaluation tasks across 10+ categories, enforcing cross-modal reasoning and multi-step tool coordination for frontier AI systems.
- Contributed to the RLHF/RLAIF feedback loop training frontier AI models, pinpointing hallucination patterns, instruction-following failures, and incorrect tool-call execution.

AI Evaluation Contributor, Freelance | Handshake AI | 2026 - Present

- Authored hard, single-answer, multi-hop search questions and reference responses designed to defeat frontier models using web search, following a strict multi-rule quality rubric and a library of authoring techniques.
- Built an automated QC checker and a public workflow walkthrough that vet submissions against the rejection rules before delivery, enforcing consistency and reducing rework.

FRONTIER AI EVALUATION PROJECTS

Confidential evaluation engagements for frontier-model providers; process and methodology only, no client task content is shared.

OpenClaw Atlas (Scale AI - Outlier): Designed a golden/silver memory-trap pipeline for authoring and QC-ing agentic RL trajectories, from story and discrepancy through failure grade, silver fix, and pytest verifiers.

Blue Shell (Scale AI - Outlier): Built a two-turn golden/silver capabilities-RL pipeline across six tools where Turn 1 must fail a rubric trap and a Turn-2 nudge must pass it 100%, end to end through QC.

Lobster Safety (Scale AI - Outlier): Ran a nine-worker safety-RL pipeline with a hard PASS/STOP gate between stages, grading trajectories against an SSOT taxonomy (F1-F10 failure types, S0-S3 severity).

Project Seal (Handshake AI): Authored adversarial multi-hop search questions that defeat frontier models under an 8-rule quality standard, backed by an automated QC checker and a public workflow walkthrough.

PROJECTS

TrajLens - Agentic Trajectory Auditor | LLM Evaluation Project (Deployed) | 2026 | [Live Demo](#)

- Built and deployed an interactive agentic-trajectory auditor (Streamlit) that grades 14 recorded tool-use traces against a 9-criterion binary rubric of deterministic Python verifiers, including safety checks for buried tool errors and unauthorized side-effect actions, across a documented 9-mode failure taxonomy; also grades any user-pasted trajectory JSON live.
- Built a rubric-vs-LLM-judge meta-evaluation: the deterministic rubric matches human gold labels 14/14 while a results-only LLM-judge scores 8/14 (Cohen's kappa 0.21), fooled by polished answers and false-positive on a good run; keyless, zero-API-cost demo on public synthetic data.

JudgeLab - LLM-as-a-Judge Reliability Lab | AI Evaluation / Meta-Evaluation Project (Deployed) | 2026 | [Live Demo](#)

- Measured GPT-4-as-a-judge against 3,355 real MT-Bench human expert votes (1,232 comparable answer pairs): 87.1% decisive agreement but Cohen's kappa 0.46, a 15.8% order-swap inconsistency rate (position bias), and a quantified verbosity preference (73.1% vs humans' 67.6%).
- Computed per-category reliability showing judge quality collapses on subjective tasks (writing kappa 0.18 vs math 0.49), with a browsable gallery of real judge-vs-human disagreement cases; reproducible on public data at zero API cost.

EvalGate - CI/CD Pipeline with Model Quality Gates | MLOps / CI/CD Project (Deployed) | 2026 | [Live Dashboard](#)

- Built a GitHub Actions CI/CD pipeline that blocks bad models: every push runs 11 pytest checks (data contracts, leakage guards, model determinism), retrains the model, and enforces a hard quality gate (test ROC-AUC ≥ 0.82 plus a CV-vs-test drift check) that fails the build if quality regresses.
- Chained CI to CD via workflow_run with artifact promotion (CD ships exactly the gated model, never retrains), auto-deploying a model-quality dashboard with deployment history to GitHub Pages, with weekly scheduled retrains.

Customer Churn Predictor | Machine Learning Project (Deployed) | 2026 | [Live Demo](#)

- Built and deployed an end-to-end churn-prediction web app: a Gradient Boosting scikit-learn Pipeline (scaling + one-hot encoding) on the 7,043-customer IBM Telco dataset, reaching 0.843 test ROC-AUC (0.847 $\pm$ 0.013 across 5-fold cross-validation); served live on Streamlit Community Cloud.
- Engineered per-prediction explainability with SHAP (one-hot contributions aggregated back to original features) and mapped risk scores to business actions (retain / monitor / no-action) via a real-time risk gauge.

Diabetes Detection Using Foot Pressure | Final-Year B.E. Project, Anna University | 2025

- Engineered an end-to-end ML pipeline: 6 FSR sensors via Arduino across 30+ walking sessions at 20Hz, Butterworth filtering in MATLAB, classifiers in Python.
- Identified and resolved session-level data leakage inflating accuracy from 99% to a realistic 92.5% by implementing GroupKFold cross-validation, demonstrating production-level ML debugging skills.
- Extracted 80+ features and benchmarked 4 models, with the Keras neural network best at 92.5%.

Spam Email Classifier | NLP Project | 2024

- Developed a complete NLP classification pipeline: NLTK preprocessing, TF-IDF vectorisation (5000 features, unigrams + bigrams), reaching 97.8% accuracy with Linear SVM on 1000+ emails.

Real-Time Face Recognition | Computer Vision Project | 2024

- Implemented a real-time face detection and recognition pipeline using OpenCV (Haar Cascades + LBPH recognizer), processing live video at 15+ FPS with a 70%+ confidence threshold on a custom dataset of 50+ images per person.

EDUCATION

Bachelor of Engineering, Computer Science & Engineering (Data Science) | Jansons Institute of Technology (Anna University), Coimbatore | 2021-2025 | CGPA 8.0/10

Higher Secondary (Class XII), MPC | Sri Sai Junior College, Kanigiri, Andhra Pradesh | 2021

Secondary (Class X) | Nagarjuna Model School, Kadapa, Andhra Pradesh | 2019 | CGPA 9.8/10

CERTIFICATIONS

AWS Academy Cloud Foundations (2022) | Microsoft 365 Productivity Advanced (2022) | Ethical Hacking & Penetration Testing (2022)